

```

require("hdf5")
require("optim")
require("nn")
require("nnx")
require("cunn")
require("cutorch")
-- require("cudnn")

```

```

cmd = torch.CmdLine()

```

```

UNSEEN = -1

```

```

-- Cmd Args

```

```

cmd:option('-datafile', '', 'data file')
cmd:option('-lm', 'nn', 'classifier to use')
cmd:option('-alpha', 0.01, 'Laplace smoothing coefficient')
cmd:option('-eta', 0.01, 'learning rate')
cmd:option('-nepochs', 3, 'number of training epochs')
cmd:option('-mb', 32, 'minibatch size')
cmd:option('-k', 1, 'ratio of noise for NCE')
cmd:option('-gpu', 0, 'whether to use gpu for training')

```

```

-- Helper function which converts a tensor to a key for table lookup

```

```

function tensor_to_key(t)
    local key = ""
    local table = torch.totable(t)
    for k, v in pairs(table) do
        key = key .. tostring(v) .. ','
    end
    return string.sub(key, 1, -2) -- Remove trailing comma
end

```

```

function count_based(smoothing)

```

```

    -- Specializd function for building unigram probability distribution (no context)

```

```

    function build_unigram()
        local unigrams = {}
        for i = 1, train_y:size(1) do
            local wi = train_y[i]
            if unigrams[wi] == nil then
                unigrams[wi] = 1
            else
                unigrams[wi] = unigrams[wi] + 1
            end
        end
    end

```

```

    for wi, count in pairs(unigrams) do
        unigrams[wi] = count / train_y:size(1)
    end

```

```

end

return unigrams
end

print('Building unigram distribution...')
local unigram = build_unigram()

-- Normalized probability distribution of unigrams
function p_unigram(wi)
  p = unigram[wi]
  if p == nil then -- Never before seen unigram, go with uniform dist
    p = (1.0 / nwords)
  end
  return p
end

-- Raw frequency of unigrams in dataset
function freq_unigram(wi)
  return p_unigram(wi) * train_y:size(1)
end

-- Create co-occurrence dictionary mapping contexts to following words
function build_ngram(x, y)
  print('Building p(w_i|w_{i-n+1},...,w_{i-1})...')
  local ngram = {}
  for i = 1, x:size(1) do
    local ctx = tensor_to_key(x[i])
    local wi = y[i]
    local val = ngram[ctx]
    if val == nil then
      ngram[ctx] = {}
      ngram[ctx][wi] = 1
    else
      local innerval = ngram[ctx][wi]
      if innerval == nil then
        ngram[ctx][wi] = 1
      else
        ngram[ctx][wi] = ngram[ctx][wi] + 1
      end
    end
  end

  if i % 100000 == 0 then
    print('Processed ' .. i .. ' training examples.')
  end
end
return ngram

```

end

-- Renormalize probabilities for each ngram

function normalize_ngram(ngram)

 print('Renormalizing ngram probabilities...')

 for ctx, wis in pairs(ngram) do

 local sum = 0

 local seen = 0 -- To calculate the number of unseen wi in given ctx

 for wi, tot in pairs(wis) do

 sum = sum + tot

 seen = seen + 1

 end

 local unseen = nwords - seen

 -- Smoothing + normalization

 if smoothing == 'mle' then

 -- How to predict unseen without smoothing for mle??

 sum = sum + (unseen * (1.0 / nwords)) -- Uniform assumption

 for wi, tot in pairs(wis) do

 ngram[ctx][wi] = ngram[ctx][wi] / sum

 end

 ngram[ctx][UNSEEN] = (1.0 / nwords) / sum

 elseif smoothing == 'lap' then

 sum = sum + (nwords * alpha) -- Add probability mass for every wi (including unseen)

 for wi, tot in pairs(wis) do

 ngram[ctx][wi] = (ngram[ctx][wi] + alpha) / sum

 end

 ngram[ctx][UNSEEN] = alpha / sum

 end

 end

end

function p_ngram(ngram, ctx, wi)

 local ctxd = ngram[ctx]

 local p = 0

 if ctxd == nil then -- Never before seen context, go with unigram

 p = p_unigram(wi)

 -- p = (1 / nwords)

 else

 p = ngram[ctx][wi]

 if p == nil then -- We've seen this context before, just not this wi

 p = ngram[ctx][UNSEEN]

 end

 end

 return p

end

```

-- Count total occurrences of context and unique types within context
function lambda_for_ctx(ngram, ctx, word)
  local unique = 0
  local total = 0
  local ctx_wi_freq = 0
  if ngram[ctx] == nil then
    return 0, 0, 0, 0 -- By defn seems like should be 1, but intuitively should be 0
  else
    for wi, tot in pairs(ngram[ctx]) do
      unique = unique + 1
      total = total + tot
      if wi == word then ctx_wi_freq = tot end
    end

    local lambda = 1 - (unique / (unique + total))
    return lambda, unique, total, ctx_wi_freq
  end
end

```

```

-- Calculates interpolated probability based on occurrences of the specific
-- context/word pair, unique words in context, total words in context, and
-- the probability of the 1-order lower context
function p_wb(ctx_wi_freq, unique, total, p_lower)
  if unique + total == 0 then -- Prevent / 0 in denom
    return 0
  else
    local num = ctx_wi_freq + unique * p_lower
    local denom = unique + total
    return num / denom
  end
end

```

```

function interpolated_p_wb(trigram, bigram, x, wi)
  if trigram ~= nil then -- Trigram, bigram, unigram interpolation
    local ctx3 = tensor_to_key(x)
    local ctx2 = tensor_to_key(torch.Tensor({x[1]}))

    local l3, unq3, tot3, ctx_wi_freq3 = lambda_for_ctx(trigram, ctx3, wi)
    local l2, unq2, tot2, ctx_wi_freq2 = lambda_for_ctx(bigram, ctx2, wi)
    local l1 = nwords / (nwords + freq_unigram(wi))
    local all = l3 + l2 + l1 -- Ensure lambdas form complex combination
    l3 = l3 / all
    l2 = l2 / all
    l1 = l1 / all

    p1 = p_unigram(wi)
  end
end

```

```

    p2 = p_wb(ctx_wi_freq2, unq2, tot2, p1)
    p3 = p_wb(ctx_wi_freq3, unq3, tot3, p2)

    return (l3 * p3) + (l2 * p2) + (l1 * p1)
else -- Bigram, unigram interpolation
    local ctx = tensor_to_key(x)

    local l2, unq2, tot2, ctx_wi_freq2 = lambda_for_ctx(bigram, ctx, wi)
    local l1 = nwords / (nwords + freq_unigram(wi))
    local all = l2 + l1 -- Ensure lambdas form complex combination
    l2 = l2 / all
    l1 = l1 / all

    p1 = p_unigram(wi)
    p2 = p_wb(ctx_wi_freq2, unq2, tot2, p1)

    return (l2 * p2) + (l1 * p1)
end
end

-- Perplexity = exponential(avg negative conditional-log-likelihood)
function perplexity(ngram, x, y)
    local sum = 0
    for i = 1, x:size(1) do
        local ctx = tensor_to_key(x[i])
        local wi = y[i]
        local p = p_ngram(ngram, ctx, wi)

        sum = sum + math.log(p)
    end

    local nll = (-1.0 / x:size(1)) * sum
    return math.exp(nll)
end

-- Perplexity = exponential(avg negative conditional-log-likelihood,
-- interpolated alongside all lower order ngrams)
function wb_perplexity(trigram, bigram, x, y)
    local sum = 0
    for i = 1, x:size(1) do
        p_interp = interpolated_p_wb(trigram, bigram, x[i], y[i])
        sum = sum + math.log(p_interp)
    end

    local nll = (-1.0 / x:size(1)) * sum
    return math.exp(nll)
end

```

```
if smoothing ~= 'wb' then
    print('Building count based model (ngram=' .. ngram .. ', smoothing=' .. smoothing ..
'...')
```

```
    local ngram = build_ngram(train_x, train_y)
    normalize_ngram(ngram)
```

```
    print('Calculating perplexity on train...')
    local perp = perplexity(ngram, train_x, train_y)
    print('Training perplexity: ' .. perp)
```

```
    print('Calculating perplexity on valid...')
    perp = perplexity(ngram, valid_x, valid_y)
    print('Validation perplexity: ' .. perp)
```

```
else
    print('Building count based model with Witten-Bell smoothing (max ngram=' ..
ngram .. ')...')
```

```
    local bigram = build_ngram(bigram_train_x, bigram_train_y)
    -- normalize_ngram(bigram)
```

```
    local trigram = nil
    if interp == 3 then
        trigram = build_ngram(trigram_train_x, trigram_train_y)
        -- normalize_ngram(trigram)
    end
```

```
    print('Calculating perplexity on train...')
    local perp = wb_perplexity(trigram, bigram, train_x, train_y)
    print('Training perplexity: ' .. perp)
```

```
    print('Calculating preplexity on valid...')
    perp = wb_perplexity(trigram, bigram, valid_x, valid_y)
    print('Validation perplexity: ' .. perp)
```

```
end
end
```

```
function nnlm(structure)
    local embedding_size = 50
    local din = embedding_size * dwin
    local dhid = 100
    local dout = nwords
```

```
    print('\nBuilding neural language model with hyperparameters:')
    print('Learning rate (eta): ' .. eta)
    print('Number of epochs (nepochs): ' .. n_epochs)
```

```

print('Mini-batch size (mb): ' .. batch_size)
print('Context window (dwin): ' .. dwin)
print('Embedding size (embed): ' .. embedding_size)
print('Vocabulary size (IVI): ' .. nwords)

-- For hierarchical softmax
function build_softmax_tree()
    hierarchy = {}
    n_buckets = math.ceil(math.sqrt(nwords))

    -- Randomly sort words into two layer tree
    for i = 1, n_buckets do
        local words = torch.range((i - 1) * n_buckets + 1, i * n_buckets):int()
        hierarchy[i] = words
    end

    return hierarchy
end

function tree_for_word(w)
    return math.ceil(w / n_buckets)
end

local model = nn.Sequential()

if structure == 'nn' then
    -- Lookup table concatenates embeddings for words in context
    input_embedding = nn.LookupTable(nwords, embedding_size)
    -- input_embedding.weight:cuda()
    -- print(input_embedding.weight)
    model:add(input_embedding)
    -- model:add(nn.LookupTable(nwords, embedding_size))
    -- model:add(nn.Reshape(din))
    model:add(nn.View(din))

    -- Linear, tanh, linear
    model:add(nn.Linear(din, dhid))
    model:add(nn.Tanh())
    model:add(nn.Linear(dhid, dout))

    model:add(nn.LogSoftMax())
    nll = nn.ClassNLLCriterion()
    print('Using softmax output.')
    if gpu == 1 then
        model:cuda()
        nll = nll:cuda()
        print('Using gpu accelerated training.')
    end
end

```

```

end
elseif structure == 'hsm' then
    build_softmax_tree()

    -- Parallel table to forward target directly to tree softmax module
    local para = nn.ParallelTable()
    local inp = nn.Sequential()
    input_embedding = nn.LookupTable(nwords, embedding_size)
    inp:add(input_embedding)
    inp:add(nn.Reshape(din))
    inp:add(nn.Linear(din, dhid))
    inp:add(nn.Tanh())
    inp:add(nn.Linear(dhid, dout))
    local out = nn.Identity()

    para:add(inp)
    para:add(out)
    model:add(para)

    model:add(nn.SoftMaxTree(dout, hierarchy, 1))
    nll = nn.TreeNLLCriterion()
    print('Using hierarchical softmax output.')
elseif structure == 'nce' then
    -- NB: currently incomplete
    local w1 = nn.Sequential()

    w1:add(nn.LookupTable(nwords, embedding_size))
    w1:add(nn.Reshape(din))

    -- Linear (no bias), tanh
    local lin1 = nn.Linear(din, dhid)
    lin1.bias = false
    w1:add(lin1)
    w1:add(nn.Tanh())

    local w2 = nn.Sequential()
    w2:add(nn.LookupTable(nwords, embedding_size))

    -- local lin2 = nn.Linear(dhid, dout)
    -- lin2.bias = false
    -- model:add(lin2)

    print('Using noise contrastive estimation.')
end

params, gradParams = model:getParameters()
if gpu == 1 then

```



```

    params = params:cuda()
    gradParams = gradParams:cuda()
end

function p_noise_sample()
    return k / (k + 1)
end

function p_true_sample()
    return 1 / (k + 1)
end

function train(e)
    -- Package selected dataset into minibatches
    local n_train_batches = math.floor(train_x:size(1) / batch_size) - 1

    order = torch.randperm(train_x:size(1))
    -- subsets = torch.range(1,word_order:size(1)-subset_size, subset_size)

    print('\nBeginning epoch ' .. e .. ' training: ' .. n_train_batches .. ' minibatches of size ' .. batch_size .. '.')
    for i = 1, n_train_batches do
        local batch_start = torch.random(i * batch_size + 1, (i + 1) * batch_size)
        local batch_end = math.min((batch_start + batch_size - 1), order:size(1))

        curr_batch = order[{{batch_start, batch_end}}]:long()
        local x = train_x:index(1, curr_batch)
        local y = train_y:index(1, curr_batch)

        if gpu == 1 then
            x = x:cuda()
            y = y:cuda()
        end

        -- Compute forward and backward pass (predictions + gradient updates)
        function run_minibatch(p)
            if params ~= p then
                params:copy(p)
            end

            -- Accumulate gradients from scratch each minibatch
            gradParams:zero()
            local loss = 0

            if structure == 'nn' then
                -- Forward pass
                local preds = model:forward(x)

```

```

    loss = nll:forward(preds, y)

    -- Backward pass
    local dLdpreds = nll:backward(preds, y)
    model:backward(preds, dLdpreds)
elseif structure == 'hsm' then
    -- Forward pass (requires input + target)
    local preds = model:forward({x, y})
    loss = nll:forward(preds, y)

    -- Backward pass
    local dLdpreds = nll:backward(preds, y)
    -- model:backward(preds, dLdpreds)
    model:backward(x, dLdpreds)
end

if i % 2000 == 0 then
    print('Completed ' .. i .. ' minibatches.')
    -- print('Loss after ' .. batch_end .. ' examples: ' .. loss)
end

return loss, gradParams
end

options = {
    learningRate = eta,
    learningRateDecay = 0.0001
    -- alpha = 0.95 -- For rmsprop
    -- momentum = 0.5
}

-- Use optim package for minibatch sgd
optim.sgd(run_minibatch, params, options)
-- optim.adagrad(run_minibatch, params, options)
-- optim.rmsprop(run_minibatch, params, options) -- Slower
end

-- Renormalize input embeddings after each epoch (max l2 norm = 1)
local threshold = 1
input_embedding.weight:renorm(2, 2, threshold):cuda()
end

function test(x, y)
    local preds = model:forward(x)
    local loss = nll:forward(preds, y)
    local perp = math.exp(loss)
    if structure == 'nn' then

```

```

    local max, yhat = preds:max(2)
    local correct = yhat:int():eq(y):double():mean() * 100
    return correct, loss, perp
else
    return 0, 0, perp
end
end
end

```

```

function hsm_perp(x, y)
    local preds = model:forward({x, y})
    -- Each pred is the likelihood of a leaf class (y)
    -- To generate distribution need to provide all desired
    -- classes and renormalize
    local loss = nll:forward(preds, y)
    local perp = math.exp(loss)

```

```

    return perp
end

```

```

function valid_acc()
    if structure == 'hsm' then
        return test({valid_x, valid_y}, valid_y)
    elseif gpu == 0 then
        return test(valid_x, valid_y)
    else
        return test(valid_x:cuda(), valid_y:cuda())
    end
end
end

```

```

function train_acc()
    if structure == 'hsm' then
        return test({train_x, train_y}, train_y)
    else
        return test(train_x, train_y)
    end
end
end

```

```

function predict_kaggle()
    local fl = torch.DiskFile('training_output/predictions_model=' .. lm .. ',dataset=' ..
datafile .. '.txt', 'w')
    fl:writeString('ID') -- Header row
    for i = 1, test_dist[1]:size(1) do
        fl:writeString(',Class' .. i)
    end
end

```

```

    local preds = model:forward(test_x)
    for i = 1, preds:size(1) do

```

```

    local targets = test_dist[i]
    local i_preds = torch.Tensor(targets:size(1))
    for j = 1, targets:size(1) do
        i_preds[j] = preds[i][targets[j]]
    end
    i_preds = i_preds / i_preds:sum()

    fl:writeString(tostring(i))
    for j = 1, i_preds:size(1) do
        fl:writeString(', ' .. i_preds[j])
    end
    fl:writeString('\n')
end

fl:close()
end

local inita, initl, initp = valid_acc()
print('Validation perplexity before training: ' .. initp .. '.')
print('Beginning training...')
local vloss = torch.DoubleTensor(n_epochs) -- valid/train loss/accuracy/time
local tloss = torch.DoubleTensor(n_epochs)
local vacc = torch.DoubleTensor(n_epochs)
local tacc = torch.DoubleTensor(n_epochs)
local etime = torch.DoubleTensor(n_epochs)

for i = 1, n_epochs do
    local timer = torch.Timer()
    train(i)

    local va, vl, vp = valid_acc()
    -- local ta, tl, tp = train_acc()

    -- vloss[i] = vl
    -- tloss[i] = tl
    vacc[i] = vp
    -- tacc[i] = tp
    etime[i] = timer:time().real

    print('Epoch ' .. i .. ' training completed in ' .. timer:time().real .. ' seconds.')
    print('Validation perplexity after epoch ' .. i .. ': ' .. vp .. '.')

    -- Logging
    local flr = torch.DiskFile('training_output/epoch=' .. i .. ',model=' .. lm .. ',dataset=' ..
datafile .. '.txt', 'w')
    for j = 1, i do
        -- flr:writeString(vacc[j] .. ' ' .. tacc[j] .. ' ' .. vloss[j] .. ' ' .. tloss[j] .. ' ' .. etime[j] .. '\n')
    end
end

```

```

        -- flr:writeString(tacc[j] .. ',' .. vacc[j] .. ',' .. etime[j] .. '\n')
        flr:writeString(vacc[j] .. ',' .. etime[j] .. '\n')
    end
    flr:close()
end

-- Kaggle predictions
predict_kaggle()

-- Export lookup table weights
local f = hdf5.open('embed_export50.hdf5', 'w')
f:write('/embed', input_embedding.weight)
f:close()

local ta, tl, tp = train_acc()
print('Final training perplexity after all epochs: ' .. tp)
end

function main()
    -- Parse input params
    opt = cmd:parse(arg)
    datafile = opt.datafile
    lm = opt.lm
    alpha = opt.alpha
    eta = opt.eta
    n_epochs = opt.nepochs
    batch_size = opt.mb
    k = opt.k
    gpu = opt.gpu

    local f = hdf5.open(opt.datafile, 'r')
    nwords = f:read('nwords'):all():long()[1]
    nclasses = f:read('nclasses'):all():long()[1]
    ngram = f:read('ngram'):all():long()[1]
    dwin = ngram - 1

    -- Split training and validation data
    train_x = f:read('train_input'):all()
    train_y = f:read('train_output'):all()
    valid_x = f:read('valid_input'):all()
    valid_y = f:read('valid_output'):all()
    test_x = f:read('test_input'):all()
    test_dist = f:read('test_output'):all()

    -- Format datasets for Witten-Bell
    if lm == 'wb' then
        if string.find(opt.datafile, '3') then

```

```

interp = 3
-- Need to load bigrams
local f2 = hdf5.open('PTB_2gram.hdf5', 'r')
bigram_train_x = f2:read('train_input'):all()
bigram_train_y = f2:read('train_output'):all()
bigram_valid_x = f2:read('valid_input'):all()
bigram_valid_y = f2:read('valid_output'):all()

trigram_train_x = train_x
trigram_train_y = train_y
trigram_valid_x = valid_x
trigram_valid_y = valid_y
else
    interp = 2
    bigram_train_x = train_x
    bigram_train_y = train_y
    bigram_valid_x = valid_x
    bigram_valid_y = valid_y
end
end
end

if lm == 'mle' or lm == 'lap' or lm == 'wb' then
    count_based(lm)
else
    nnlm(lm)
end
end

main()

```